

Parallel Computing with GPUs: Warp Level CUDA micro

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Dr Paul Ric

<http://paulrichmond.shef.ac.uk/> Add WeChat [edu_assist_pro](#) COM4521/



The
University
Of
Sheffield.

 NVIDIA

GPU
RESEARCH
CENTER

Last Teaching Week

- ❑ We learnt about shared memory
 - ❑ Very powerful for block level computations
 - ❑ Excellent for improving performance by reducing memory bandwidth
 - ❑ User controlled caching and needs careful consideration for bank conflicts and boundary conditions
- ❑ Memory coalescing: V bandwidth performance
 - ❑ Need to be aware of c
- ❑ Occupancy can be changed by modifying k-sizes, registers and shared memory usage
- ❑ This week:
 - ❑ How exactly are warps scheduled?
 - ❑ Can we program at the warp level?
 - ❑ What mechanisms are there for communication between threads?

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Overview

❑ Warp Scheduling & Divergence

❑ Atomics

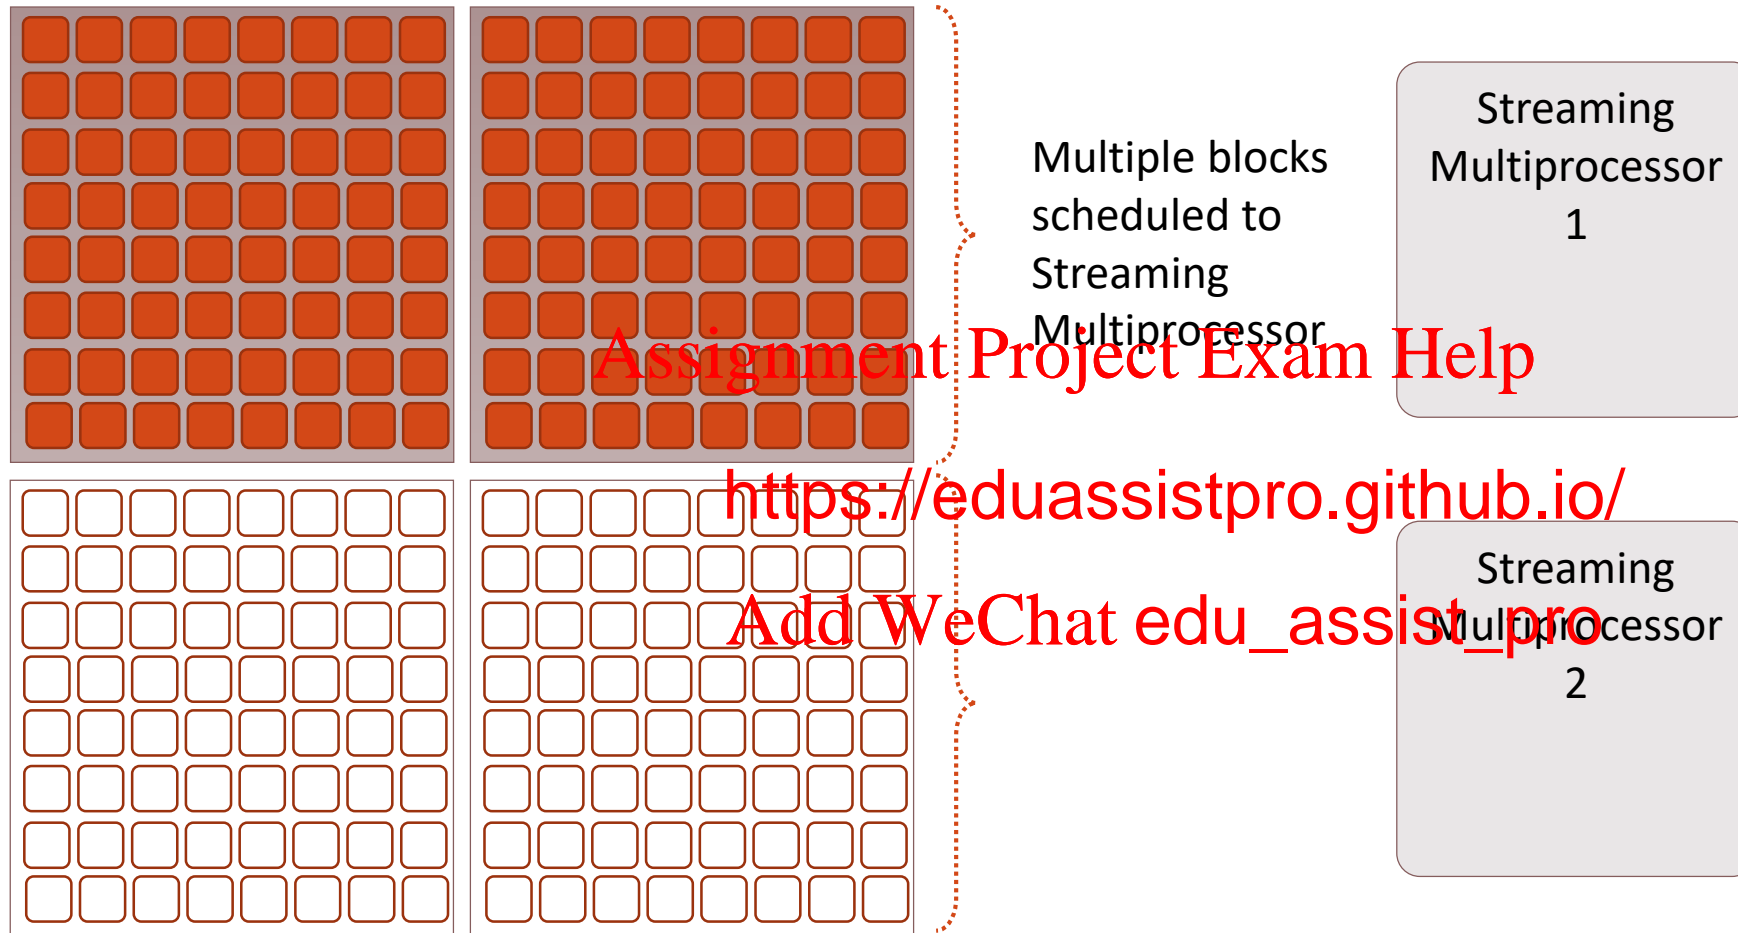
❑ Warp Operations

Assignment Project Exam Help

<https://eduassistpro.github.io/>

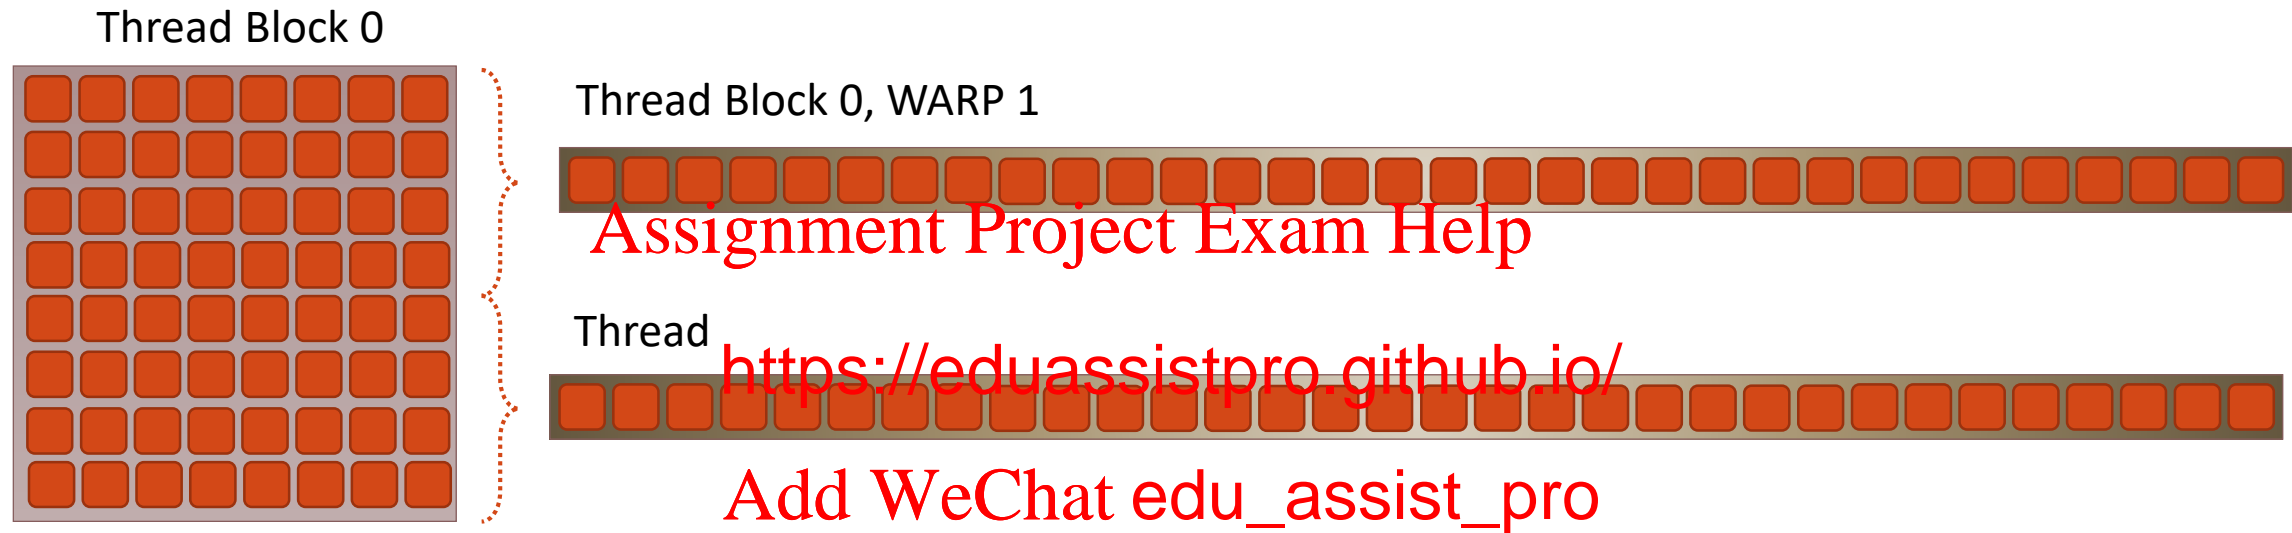
Add WeChat edu_assist_pro

Thread Block Scheduling



- ❑ No guarantee of block ordering on SMPs
- ❑ Hardware will schedule blocks to a SMP as soon as necessary resources are available

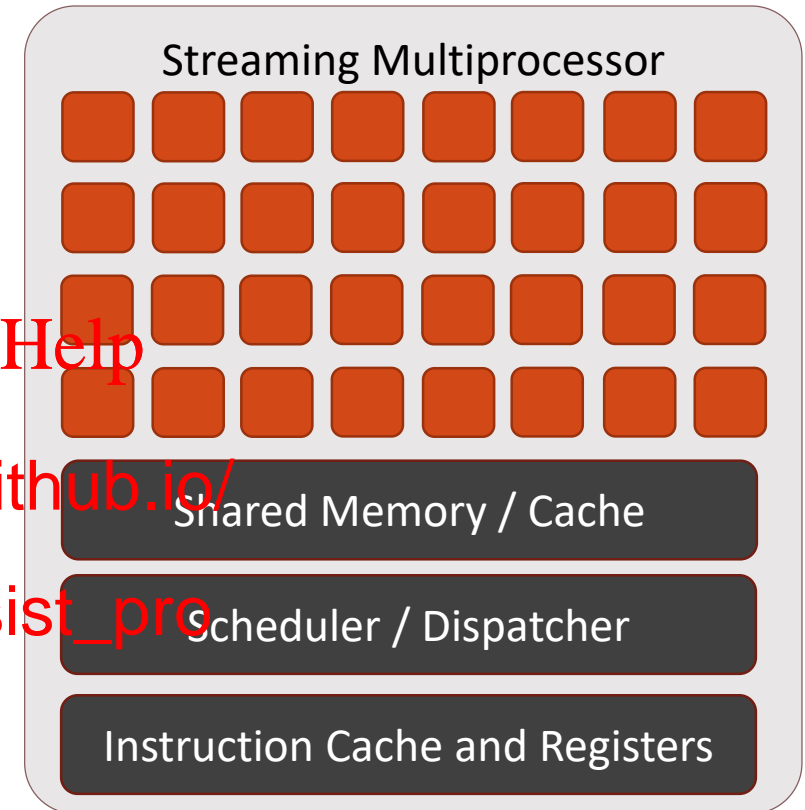
Thread Block Scheduling



- ❑ Each thread block is mapped to one or more warps
- ❑ 2D blocks are split into warps first by x index then y then z

Warp Scheduling

- ❑ Zero overhead to swap warps (warp scheduling)
 - ❑ Warps contain only threads from a single thread block
 - ❑ Warps can be swapped between different blocks assigned to a streaming multiprocessor
 - ❑ At any one time only one warp has operations being executed
 - ❑ Memory movement happens in background



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Warps and SIMD

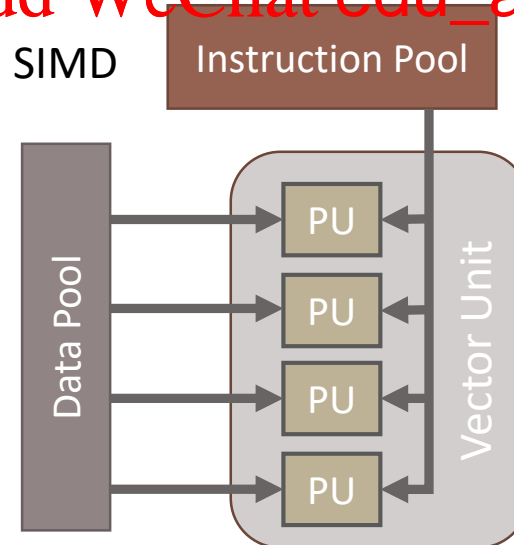
- ❑ Execution of GPU instructions is always in groups of threads called **warps**
- ❑ Within a warp execution on the hardware follows the SIMD execution model
 - ❑ The view outside of a
- ❑ What happens if code
 - ❑ **Branch Divergence**

Assignment Project Exam Help

<https://eduassistpro.github.io/>

erent control flow?

Add WeChat edu_assist_pro



Divergent Threads

- ❑ All threads must follow SIMD model

 - ❑ Multiple code branch paths must be evaluated

 - ❑ Not all threads will be active during code execution

 - ❑ Coherence = all threads following the same path

Assignment Project Exam Help

- ❑ How to avoid diverge <https://eduassistpro.github.io/>

1. Avoid conditional code

Add WeChat edu_assist_pro

2. **Especially** avoid conditional code based on `adIdx`

- ❑ Fully coherent code can still have branches

 - ❑ BUT all threads in the warp follow the same path



Coherent Code

```
__global__ void a_kernel()  
{  
    if (blockIdx.x % 2)  
        //something  
    else  
        //something else  
}
```

```
__global__ void b_kernel()  
{  
    if (threadIdx.x % 2)  
        //something  
    e  
        //something else  
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

☐ Which is coherent?

☐ Which is divergent?

Levels of divergence

- ❑ Divergent code can be classified by how many “ways” it diverges.
 - ❑ E.g. the following examples are 4-way divergent (and functionally equivalent)
- ❑ If a warp has 32-way divergence this will have a **massive** impact on performance!

Assignment Project Exam Help

<https://eduassistpro.github.io/>

```
__global__ void a_kernel(int *a)
{
    int a = a[threadIdx.x + blockIdx.x*blockDim.x]
    if (a==0)
        //code for case 0
    else if (a==1)
        //code for case 1
    else if (a==2)
        //code for case 2
    else if (a==3)
        //code for case 3
}
```


Add WeChat [edu_assist_pro](#)

```
void a_kernel(int *a)
{
    int a = a[threadIdx.x + blockIdx.x*blockDim.x]
    switch (a){
        case(0):
            //code for case 0 with break
        case(1):
            //code for case 1 with break
        case(2):
            //code for case 2 with break
        case(3):
            //code for case 3 with break
    }
}
```



2D blocks and divergence

Thread Block 0



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
__global__ void a_kernel()  
{  
    if (threadIdx.y % 2)  
        //something  
    else  
        //something else  
}
```

```
__global__ void b_kernel()  
{  
    if (threadIdx.y / 4)  
        //something  
    else  
        //something else  
}
```

❑ How many ways of divergence?

Branching vs. Predication

- ❑ Predication is an optional guard that can be applied to machine instructions
 - ❑ A predicate is set in predicate registers (virtual registers)
 - ❑ Predicates are unique to each thread
- ❑ Depending on the pre-condition, execution can be conditionally executed
 - ❑ NOP otherwise
- ❑ How does this differ to branching?
 - ❑ No labels or change in program counter
 - ❑ Smaller more compact code
 - ❑ Less operations = better performance

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Branching code

CUDA C

```
int a = 0;

if (i < n)
    a = 1;
else
    a = 2;
```

PTX ISA

```
mov.s32 a, 0;
setp.lt.s32 p, i, n;
@!p bra A_FALSE;
A_TRUE: //if true
        mov.s32 a, 1; //a=1
        bra A_END;
A_FALSE: //if false
        mov.s32 a, 2; //a=2
A_END:
...
```

❑ Consider the following branching code...

❑ Code is PTX ISA

low-level parallel thread execution virtual machine and instruction set architecture

for CUDA

dependent of NVIDIA GPU

architecture

❑ Used to generate native target architecture machine instructions

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

Branching code using predicate

CUDA C

```
int a = 0;

if (i < n)
    a = 1;
else
    a = 2;
```

PTX ISA (compiler optimised)

```
mov.s32 a, 0;
setp.lt.s32 p, i, n; //p=(i<n)
@p mov.s32 a, 1; //a=1
@!p mov.s32 a, 2; //a=2
```

CUDA C (improved)

```
int a = 0;
a = (i < n)? 1: 2;
```

PTX ISA (improved)

```
mov.s32 a, 0; //a=0
setp.lt.s32 p, i, n; //p=(i<n)
selp a, 1, 2, p //a=(p)?1:2
```

❑ Consider the following branching code...

❑ In this case the predicate can be used to reduce the number of instructions

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

good at

branching and predication

❑ Can hint to the compiler by using ternary operators

❑ Warp Scheduling & Divergence

❑ Atomics

❑ Warp Operations

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



What is wrong with the following

Assignment Project Exam Help

```
__global__ void m
{
    __shared__ int
    int my_local = a[threadIdx.x] + lock[m.x];

    if (my_local > max)
        max = my_local;
}
```

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Atomics

- ❑ Atomics are used to ensure correctness when concurrently reading and writing to a memory location (global or shared)

Assignment Project Exam Help

```
__global__ void m
{
    __shared__ int
    int my_local = threadIdx.x + blockDim.x;

    if (my_local > max)
        max = atomicMax(&max, my_local);
}
```

<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

- ❑ No race condition
- ❑ Function *supported* in (some) hardware
 - ❑ Support varies depending on which memory is used (global, shared etc.)

Atomic Functions and Locks

❑ An atomic function

- ❑ Must guarantee that an operation can complete without interference from any other thread
- ❑ Does not provide any guarantee of ordering or provide any synchronisation

❑ How can we impleme

```
__device__ int lock = 0;

__global__ void kernel() {
    bool need_lock = true;
    // get lock
    while (need_lock) {
        if (atomicCAS(&lock, 0, 1)==0) {
            //critical code section
            atomicExch(&lock, 0);
            need_lock = false;
        }
    }
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

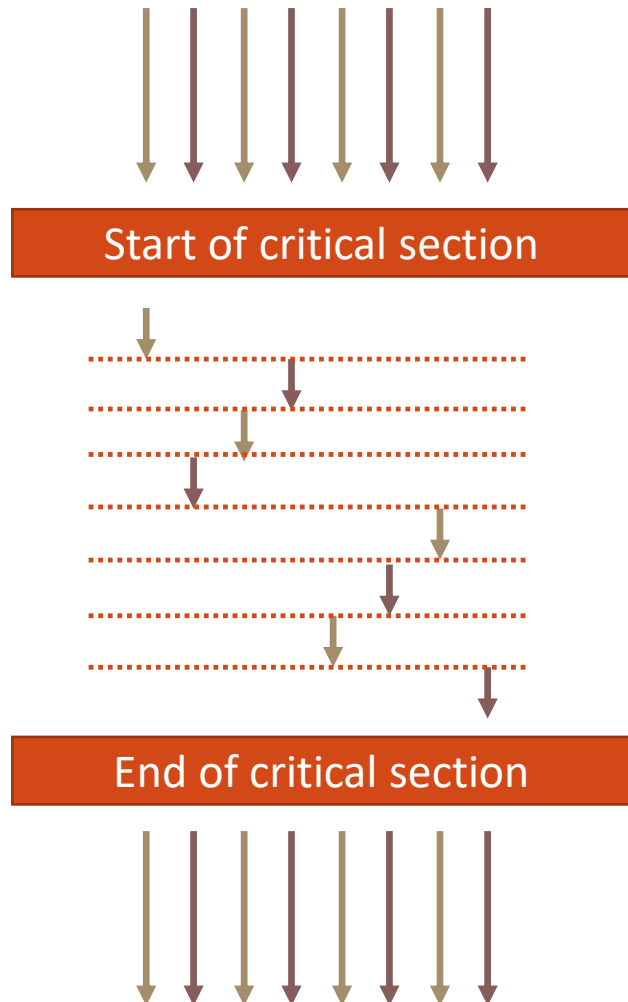
int ato address, int compare, int val)

Performs the following in a single atomic transaction (atomic instruction)

*address=(*address==compare)? val : *address;

Returning the old value at the address

Serialisation



❑ What happens to performance when using atomics?

Assignment Project Exam Help

critical section example

<https://eduassistpro.github.io/>

for each thread accessing

the sh

Add WeChat edu_assist_pro

❑ For the

instruction access to

the shared lock variable is serialised

❑ This is true of any atomic function or instruction in CUDA

CUDA Atomic Functions / Instructions

❑ In addition to `atomicCAS` the following atomic functions/instructions are available

❑ Addition/subtraction

❑ E.g. `int atomicAdd(int* address, int val)` – add `val` to integer at `address`

❑ Exchange

❑ Exchange a value with <https://eduassistpro.github.io/>

❑ Increment/Decrement

❑ Minimum and Maximum

❑ Variants of atomic functions

❑ Floating point versions require Compute 2.0

❑ 64 bit integer and double versions available in Pascal (Compute 6.0)

❑ See docs: <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomic-functions>

Assignment Project Exam Help

Add WeChat edu_assist_pro

Shared vs Global Atomics

❑ Global Atomics

- ❑ Fermi: Atomics are not cached and are hence very slow
- ❑ Kepler and Maxwell: both use L2 caching of global atomics

❑ Shared Memory Atomics

- ❑ Fermi and Kepler: No atomics
 - ❑ Emulated using locks i
 - ❑ Poor when there is high contention
 - ❑ Sometimes worse than global atomics
- ❑ Maxwell+: Hardware supported SM atomics
 - ❑ Much improved performance

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Local vs Global Atomics

Kepler

Maxwell

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

☐ Image histogram example

- ☐ Accumulation of colour values for images

- ☐ Entropy: measure of the level of disorder (lower entropy == higher contention)

- ☐ <https://devblogs.nvidia.com/parallelforall/gpu-pro-tip-fast-histograms-using-shared-atomics-maxwell/>

❑ Warp Scheduling & Divergence

❑ Atomics

❑ Warp Operations

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Warp Shuffle

- ☐ For moving/comparing data between threads in a block it is possible to use Shared Memory (SM)
- ☐ For moving/comparing data between threads in a warp (known as lanes in this context) it is possible to use a *warp shuffle* (SHFL)
 - ☐ Direct exchange of information between two threads
 - ☐ Can replace atomics
 - ☐ Should never depend
 - ☐ Does not require SM
 - ☐ Always faster than SM equivalent
 - ☐ Implicit synchronisation (no need for `__syncthreads`)
 - ☐ EXCEPT on Volta hardware
 - ☐ Works by allowing threads to read another threads registers
 - ☐ Available on Kepler and Maxwell


Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Shuffle Variants

`__shfl()`



Shuffled between
any two index
threads

`__shfl_up()`




Shuffles to n^{th} right
neighbour wrapping
indices (in this case
 $n=2$)

`__shfl_down()`



n^{th} left
neighbour wrapping
is case
 $n=2$)

`__shfl_xor()`



Butterfly (XOR)
exchange shuffle
pattern

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Shuffle function arguments

❑ `int __shfl(int var, int srcLane, int width=warpSize);`

❑ Direct copy of var in srcLane

❑ `int __shfl_up(int var, unsigned int delta, int width=warpSize);`

❑ `int __shfl_down(int var, unsigned int delta, int width=warpSize);`

❑ delta is the n step used for shuffling

❑ `int __shfl_xor(int var, int laneMask, int width=warpSize);`

❑ Source lane determined by b

Assignment Project Exam Help

<https://eduassistpro.github.io/>

❑ Optional width argument

❑ Must be a power of 2 and less than or equal to w

❑ If smaller than warp size each subsection acts independently (own wrapping)

❑ All functions available as float and half versions

❑ <http://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#warp-shuffle-functions>

Add WeChat edu_assist_pro

Shuffle Warp Sum Example (down)

```
__global__ void sum_warp_kernel_shfl_down(int *a)
{
    int local_sum = a[threadIdx.x + blockIdx.x*blockDim.x];


    for (int offset = WARP_SIZE / 2; offset>0; offset /= 2)
        local_sum += __shfl_down(local_sum, offset);

    if (threadIdx.x%32 == 0)
        printf("Warp max is %d", local_sum);
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Shuffle Warp Sum Example (xor)

```
__global__ void sum_warp_kernel_shfl_xor(int *a)
{
    int local_sum = a[threadIdx.x + blockIdx.x*blockDim.x];

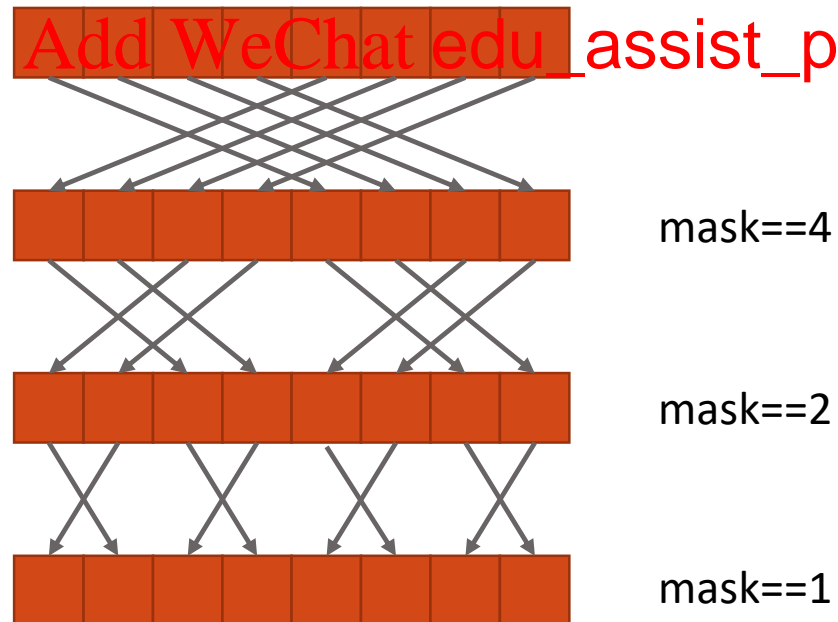
    for (int mask = WARP_SIZE / 2; mask>0; mask /= 2)
        local_sum += __shfl_xor(local_sum, mask);

    if (threadIdx.x%32 == 0)
        printf("Warp max is %d", local_sum);
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Warp sum in all threads

Warp Voting

- ❑ Warp shuffles allow data to be exchanged between threads in a warp
- ❑ Warp voting allows threads to test a condition across all threads in a warp
 - ❑ `int all(condition)`
 - ❑ True if the condition is
 - ❑ `int any(condition)`
 - ❑ True if any thread in warp meets condition
 - ❑ `unsigned int ballot(condition)`
 - ❑ Sets the n^{th} bit of the return value based on the n^{th} threads condition value
- ❑ All warp voting functions are single instruction and act as barrier
 - ❑ Only active threads participate, does not block like `syncthreads()`

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat `edu_assist_pro`

Warp Voting Example

```
__global__ void voteAllKernel(unsigned int *input, unsigned int *result)
{
    int i = threadIdx.x + blockIdx.x*blockDim.x;
    int j = i % WARP_SIZE;

    int vote_result = all(i

    if (j==0)
        result[j] = vote_result;
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- ❑ For each first thread in the warp calculate if all threads in the warp have `true` valued input
- ❑ Save the warp vote to a compact array
 - ❑ A reduction of factor 32

__shfl_sync

❑ Volta hardware allows interleaved execution of statements from divergent branches

❑ Each thread has its own program counter to allow this

Pre-Volta hardware ❑ As a result warp operations require a serialised version

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro


DA 8:

```
shfl(int var, int srcLane, int  
pSize);
```

9:

```
hfl(unsigned int mask, int var, int srcLane,  
int width=warpSize);
```

```
if (threadIdx.x < 4) {  
    A;  
    B;  
} else {  
    X;  
    Y;  
}  
Z;  
__syncwarp();
```



❑ A mask of `0xFFFFFFFF` will sync whole warp and act like CUDA 8 shuffle

❑ Allow syncing of units smaller than a warp

Global Communication

- ❑ Shared memory is per thread block
- ❑ Shuffles and voting for warp level
- ❑ Atomics can be used for some global (grid wide) operations
- ❑ What about general g
 - ❑ Not possible within a <https://eduassistpro.github.io/>!
 - ❑ Remember a grid may not be entirely the device
 - ❑ Can be enforced by finishing the kern

```
step1 <<<grid, blk >>>(input, step1_output);  
// step1_output can safely be used as input for step2  
step2 <<<grid, blk >>>(step1_output, step2_output);
```


Summary

- ❑ Warps are the level in which threads are grouped for execution
- ❑ Divergent code paths within a warp are very bad for performance
- ❑ Warps can communicate directly via warp shuffles and voting
- ❑ The performance of w very fast (single instruction)
<https://eduassistpro.github.io/>
- ❑ Atomic can be used to allow thread tive access to a shared variable
Add WeChat edu_assist_pro
- ❑ Atomic performance varies greatly with different architectures

Acknowledgements and Further Reading

- ❑ Predication: <http://docs.nvidia.com/cuda/parallel-thread-execution/index.html#predicated-execution>
 - ❑ Shuffling: <http://on-demand.gputechconf.com/gtc/2013/presentation/S3174-Kepler-Shuffle-Tips-Tricks.pdf>
 - ❑ Volta: <https://devblogs.nvidia.com/parallelforall/volta-architecture-revealed/>
- <https://eduassistpro.github.io/>
Add WeChat edu_assist_pro