

Parallel Computing with GPUs: Sorting an ies

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Dr Paul Ric

<http://paulrichmond.shef.ac.uk/> Add WeChat [edu_assist_pro](https://eduassistpro.github.io/) COM4521/



The
University
Of
Sheffield.

 NVIDIA

GPU
RESEARCH
CENTER

Last Week

- ❑ We learnt about Performance optimisation
- ❑ APOD cycle
- ❑ Use of guided analysis to find important kernels
- ❑ Use of guided analysis

Assignment Project Exam Help

outes for code

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Important Reminder

☐ Guest lecture next week

☐ MOLE Quiz next week 9.00am

☐ Followed by 1 hour lab (assignment help and lab catchup)

☐ Week 11: **Assignment Project Exam Help**

☐ No lecture (bank holiday)

☐ Lab for assignment help

<https://eduassistpro.github.io/>

☐ Week 12:

Add WeChat edu_assist_pro

☐ No lecture or lab

❑ Sorting Networks

❑ Merge and Bitonic sort

❑ Thrust Parallel Primitives Library

❑ Applications of sorting

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Serial Sorting Examples

❑ Insertion Sort

❑ Insert a new element into a sorted list.

❑ E.g. [1 6 3 4 2 5]

❑ [1] -> [1 6] -> [1 3 6] -> [1 3 4 6] -> [1 2 3 4 6] -> [1 2 3 4 5 6]

❑ Bubble Sort

❑ Exchange and Sweep through adjacent elements

❑ $O(n^2)$ worst-case and average case, $O(n)$ best case

❑ E.g. [1 6 3 4 2 5]

❑ [1 6 3 4 2 5] -> [1 3 6 4 2 5] -> [1 3 4 6 2 5] -> [1 3 4 2 6 5] -> [1 3 4 2 5 6]

❑ [1 3 2 4 5 6]

❑ [1 2 3 4 5 6]

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Classifying Sort Techniques/Implementations

☐ Data driven

- ☐ Each step of the algorithm depends on the previous step version
- ☐ Highly serial

☐ Data independent

- ☐ The algorithms perform the same operation, not change its processing based on data
- ☐ Well suited to parallel implementation
- ☐ Can be expressed as a sorting network


Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Sorting Networks

- ❑ A sorting network is a comparator network that sorts all input sequences
- ❑ Following the same execution of stages
- ❑ Consider the previous Assignment Project Exam Help



<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

[3 4 6 2 5] -> [1 3 4 2 6 5] -> [1 3 4 2 5 6]

[1 3 4 2 5 6] -> [1 3 2 4 5 6]

[1 3 2 4 5 6] -> [1 2 3 4 5 6]

[1 2 3 4 5 6]

Sweeps

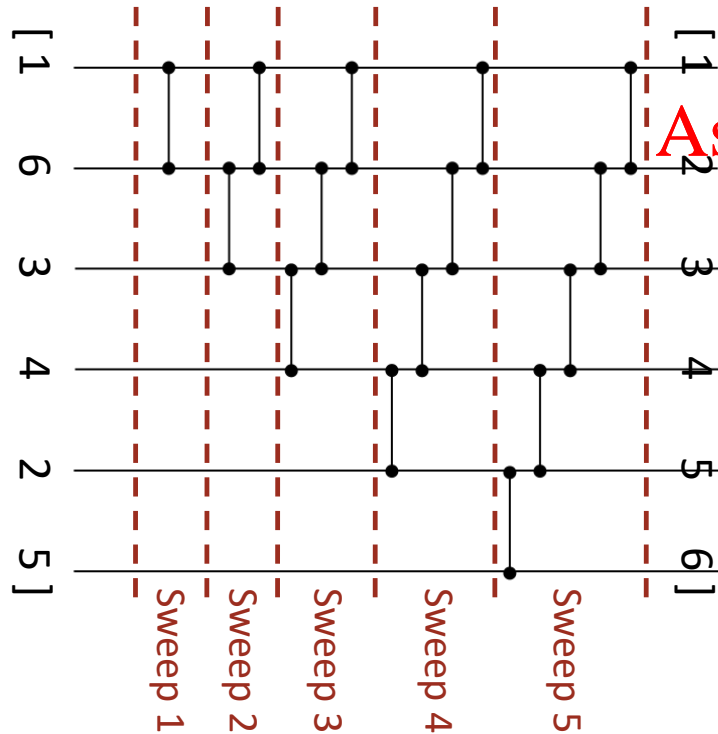
Not considered

Compared not swapped

Compared and swapped

Sorting Networks

□ And Insertion Sort...



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro


[1 6 3 4 2 5] → [1 2 3 4 5 6]
[1 6 3 4 2 5] → [1 3 6 4 2 5]
[1 3 6 4 2 5] → [1 3 4 6 2 5]
[1 3 4 6 2 5] → [1 3 4 2 6 5]
[1 3 4 2 6 5] → [1 2 3 4 6 5]
[1 2 3 4 6 5] → [1 2 3 4 5 6]
[1 2 3 4 5 6] → [1 2 3 4 5 6]

Not considered

Compared not swapped

Compared and swapped

Parallel Sorting Networks



Bubble

Insertion

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

❑ Parallel Bubble and Insertion sorting network is still not very efficient

❑ $2n - 3$ sweeps

❑ $n(n - 1)/2$ comparisons - $O(n^2)$ complexity

1 6 3 4 2 5]
1 3 4 2 5]
[1 3 4 6 2 5]
[1 3 4 2 6 5]
[1 3 2 4 5 6]
[1 2 3 4 5 6]
[1 2 3 4 5 6]
[1 2 3 4 5 6]

Sweeps = 9

❑ Sorting Networks

❑ Merge and Bitonic sort

❑ Thrust Parallel Primitives Library

❑ Applications of sorting



Assignment Project Exam Help

<https://eduassistpro.github.io/>

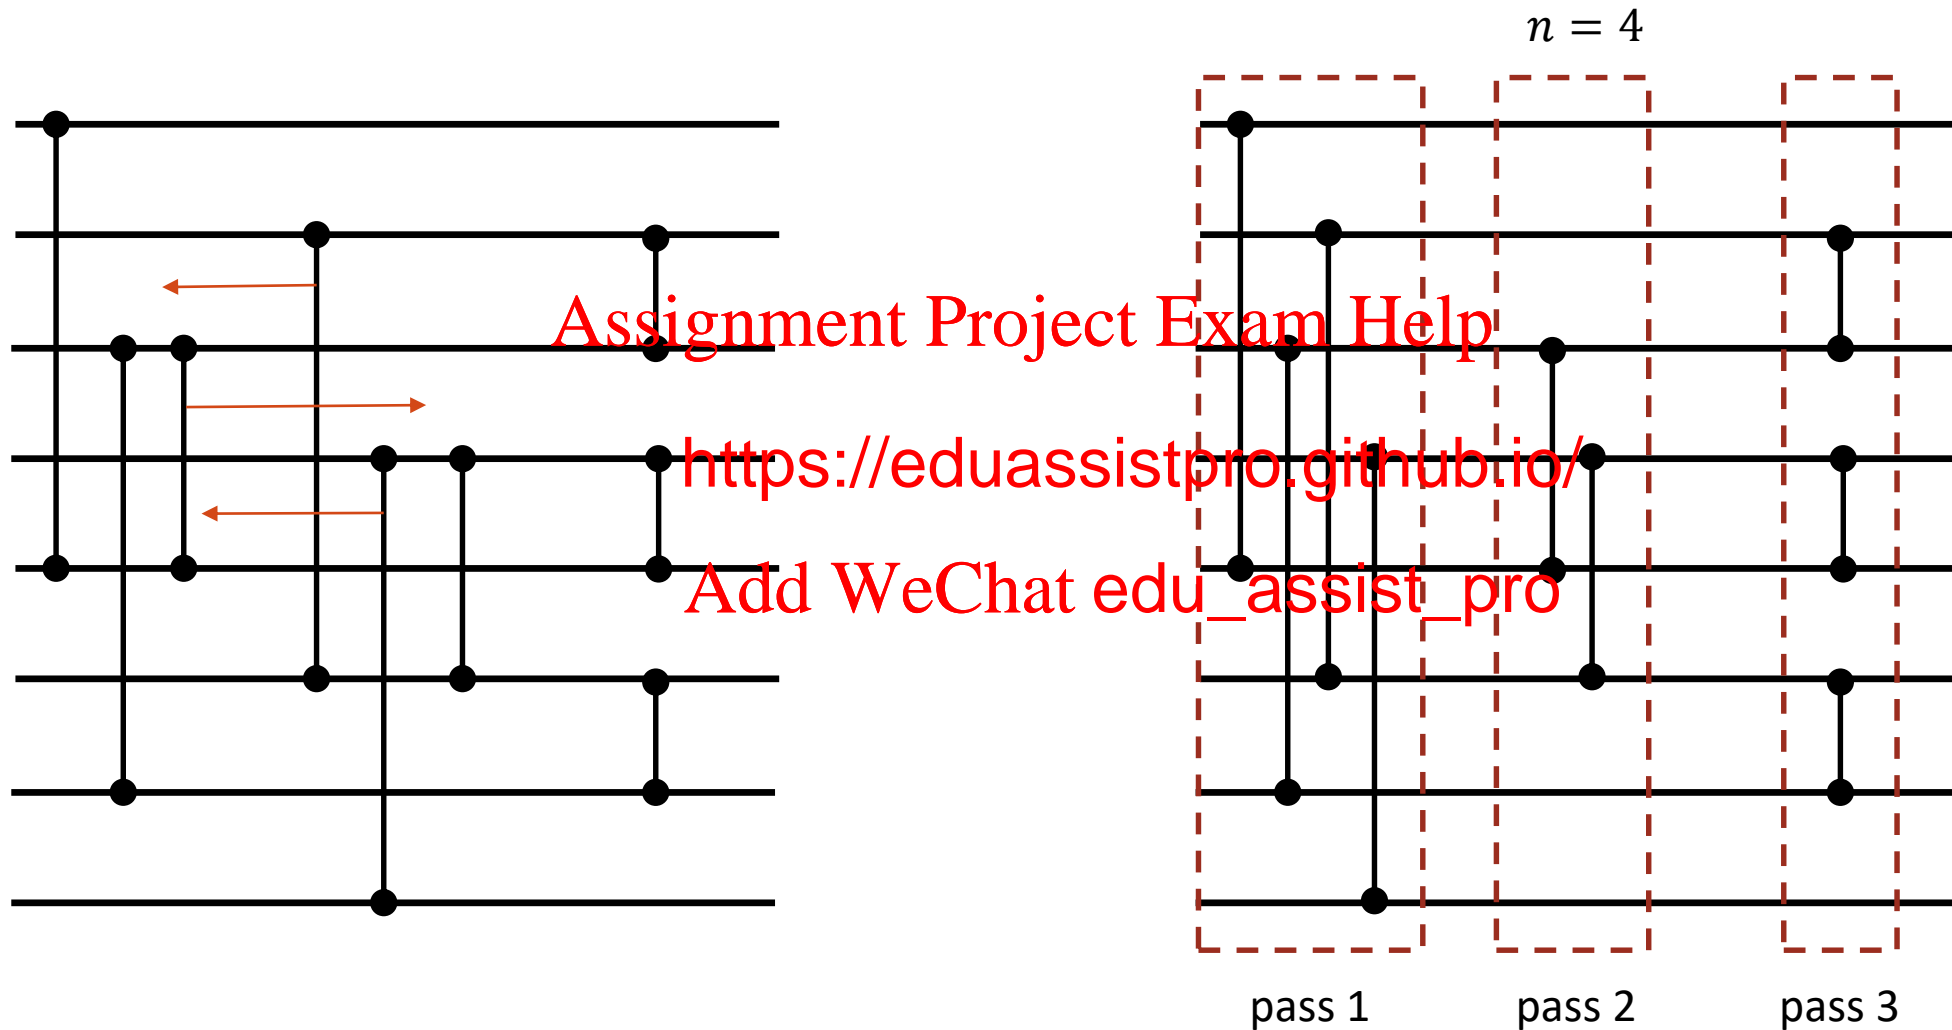
Add WeChat edu_assist_pro

Merge Sort

- ❑ To reduce the $O(n^2)$ overhead we need a better sorting network
- ❑ The odd-even merge sort network (for power of 2 n)
 - ❑ Sort all odd and even keys separately and then merge m values of a stage
 - ❑ Merge a sorted sequence of elements on lines $\langle a_1, \dots, a_n \rangle$ with those on lines $\langle a_{n+1}, \dots, a_{2n} \rangle$
 - ❑ Each merge requires 1
 - ❑ Total complexity of $O(\dots)$



Merging of two sorted sequences ($n=4$)

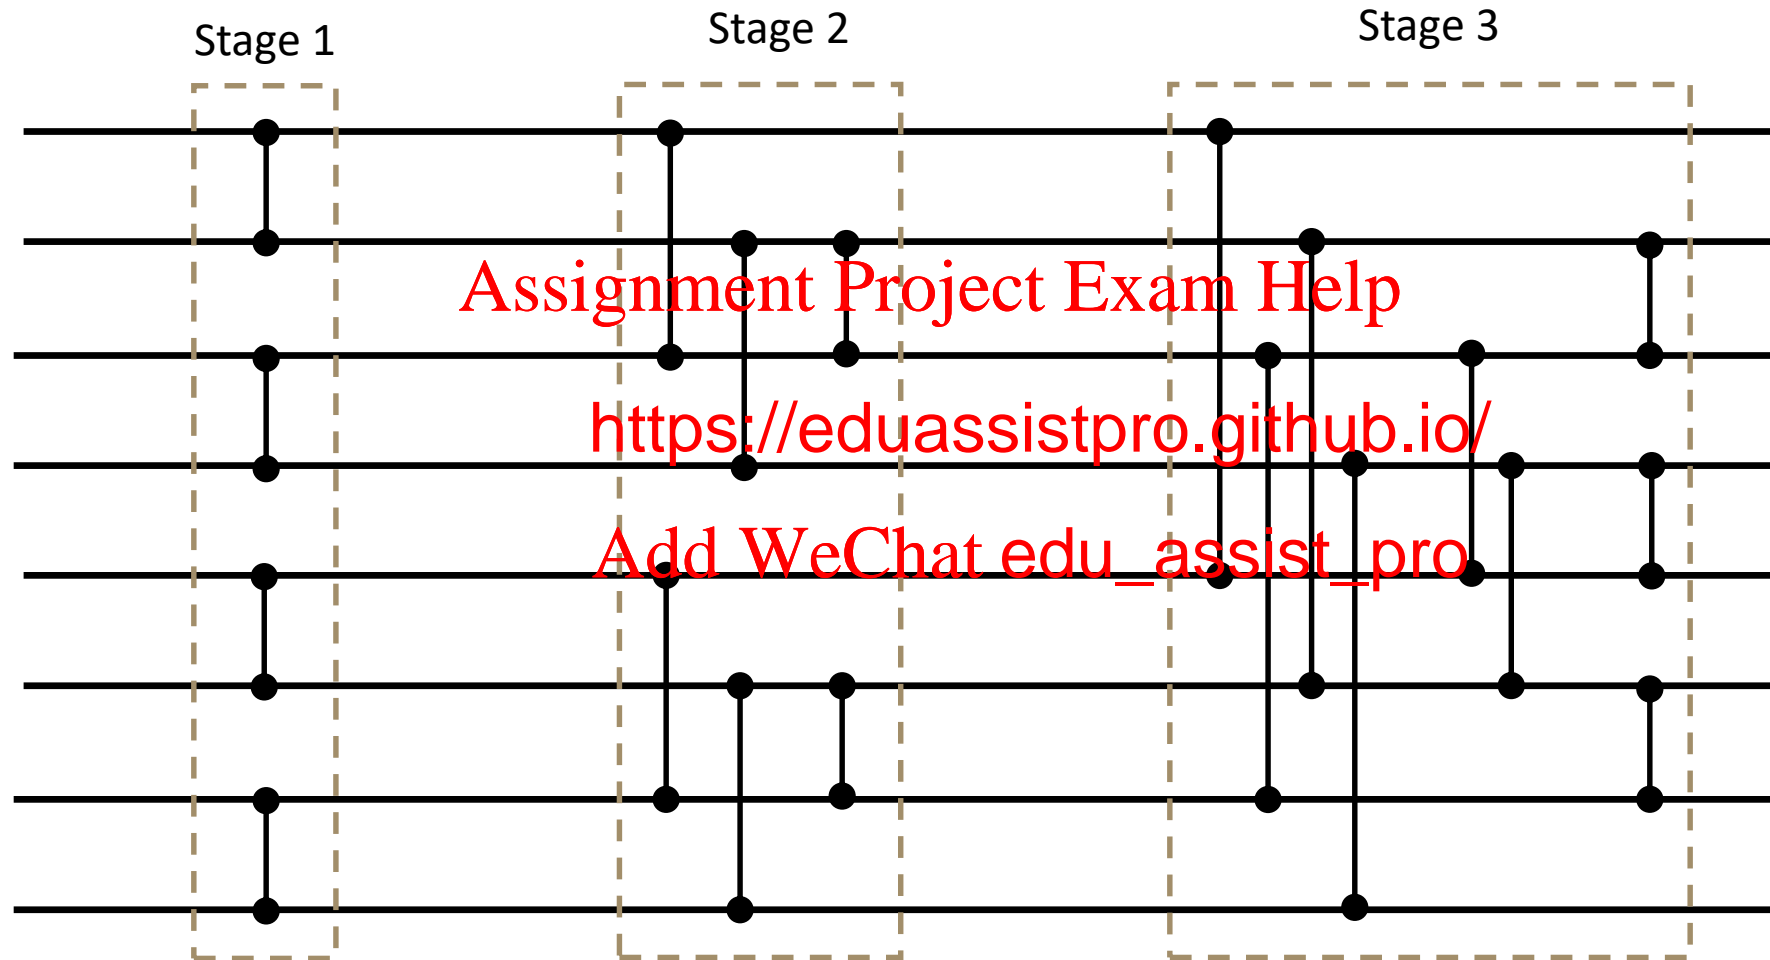


Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Merge Sorting (n=8)

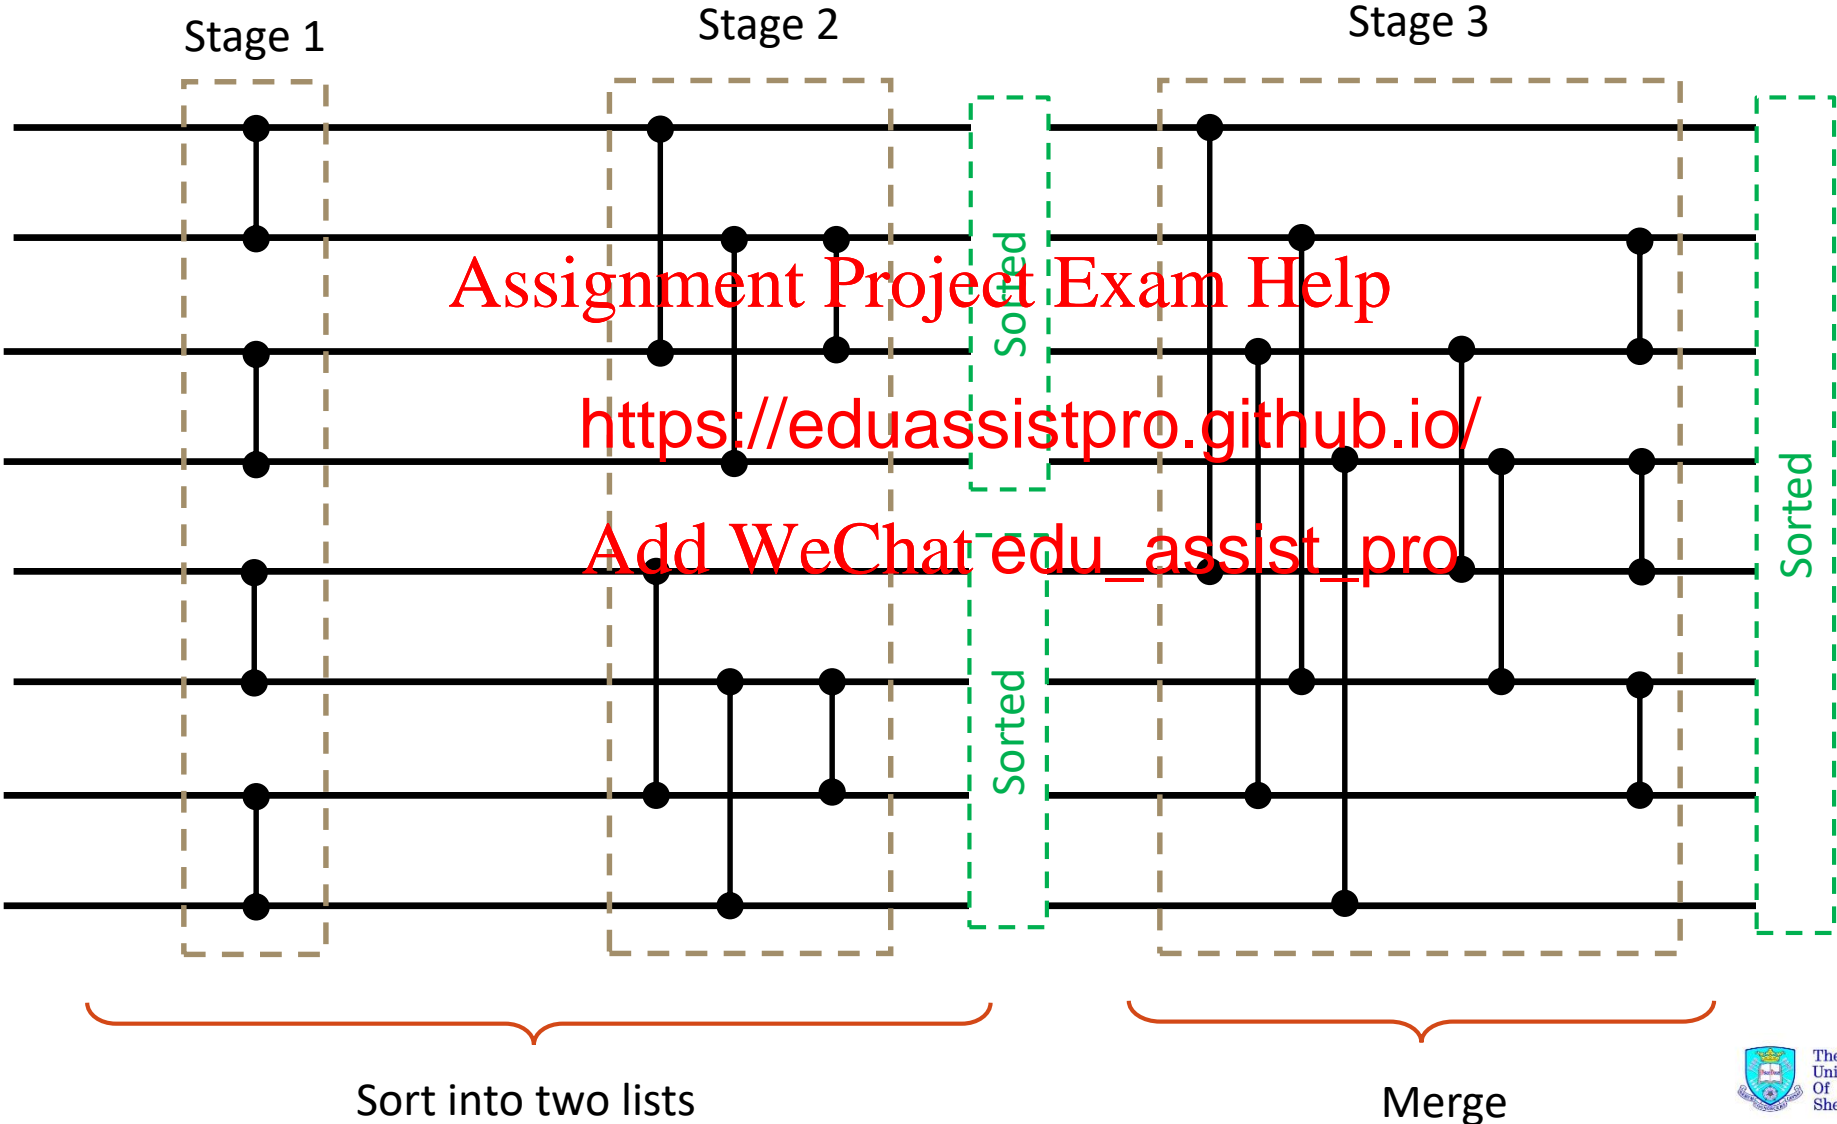


Assignment Project Exam Help


<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Merge Sorting (n=8)



Merge Sorting (n=8) example



Assignment Project Exam Help

<https://eduassistpro.github.io/>

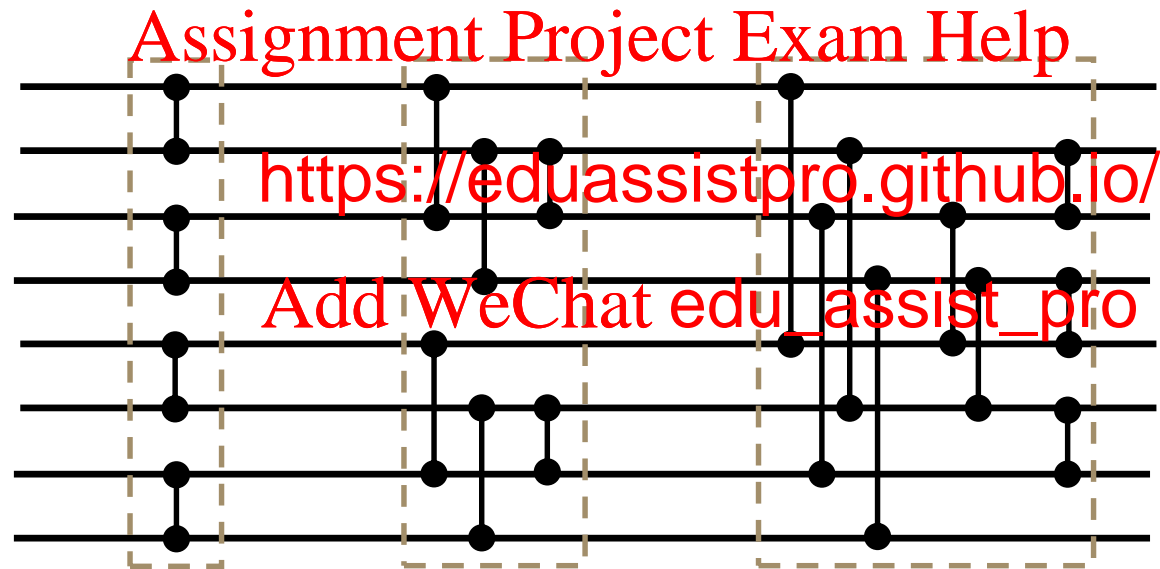
Add WeChat edu_assist_pro

Input	Stage 1		Stage 2		Stage 3						Output
8	1		1		1						1
1	8			5 3			3			2	2
5	3		3		5		5		2	3	3
3	5			8				8	4 4		4
6	2		2			2			5	5	5
2	6			6 4			4			8 6	6
4	4		4		6		6			8	8
9	9			9				9			9



Limitations of Merge Sort?

- ❑ What is potentially wrong with a merge sort GPU implementation?
 - ❑ Hint: Think about workload per thread




Limitations of Merge Sort

- ❑ What is potentially wrong with a merge sort GPU implementation?
 - ❑ Irregular memory accesses
 - ❑ Not all values are compared in each pass (uneven workload per thread)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



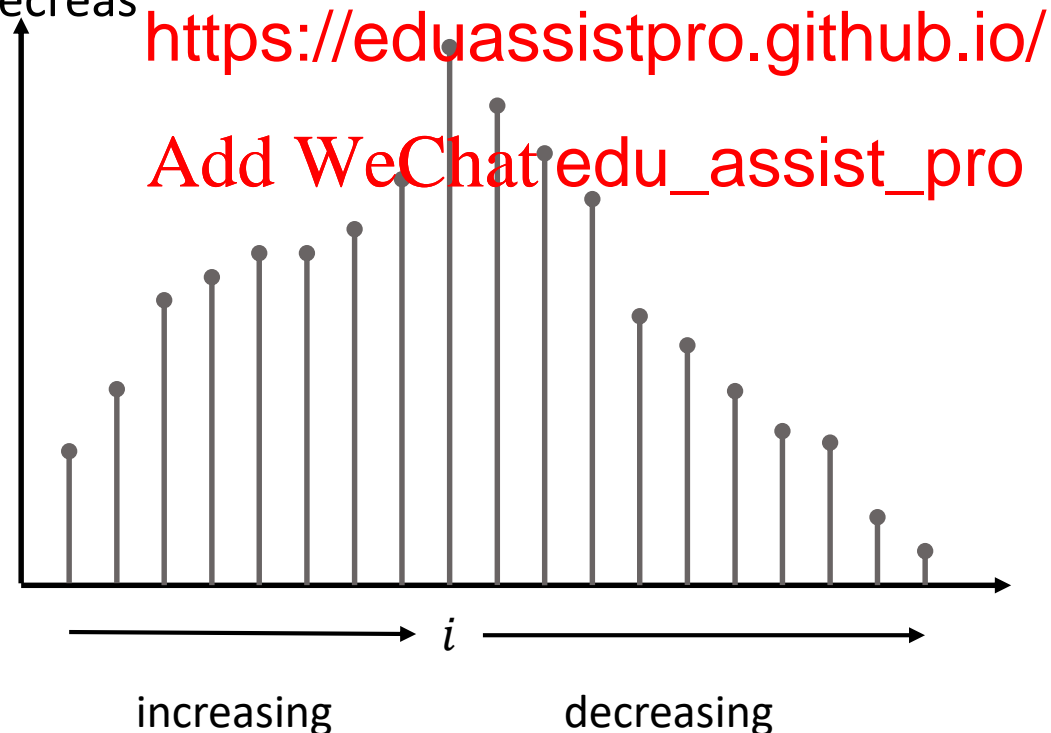
Solution: Bitonic Sort

□ Bitonic sorting network

□ Iterative splitting and merging of inputs into increasing large bitonic sequences

□ A sequence is bitonic if

□ There is an i , such that a_1, \dots, a_i is monotonically increasing and a_i, \dots, a_n is monotonically decreasing



Bitonic Sorting Network

Assignment Project Exam Help

<https://eduassistpro.github.io/>

❑ Sorting and Merging increasing large sequences

❑ When $n = 2^k$ there are k levels with $\frac{n}{2}$ comparisons each

❑ GPU Implementation

❑ Regular access strides :-)

❑ Efficiently balanced workload :-)

❑ Requires multiple kernel launches to merge over $n > \text{block size}$

❑ Sorting Networks

❑ Merge and Bitonic sort

❑ Thrust Parallel Primitives Library

❑ Applications of sorting

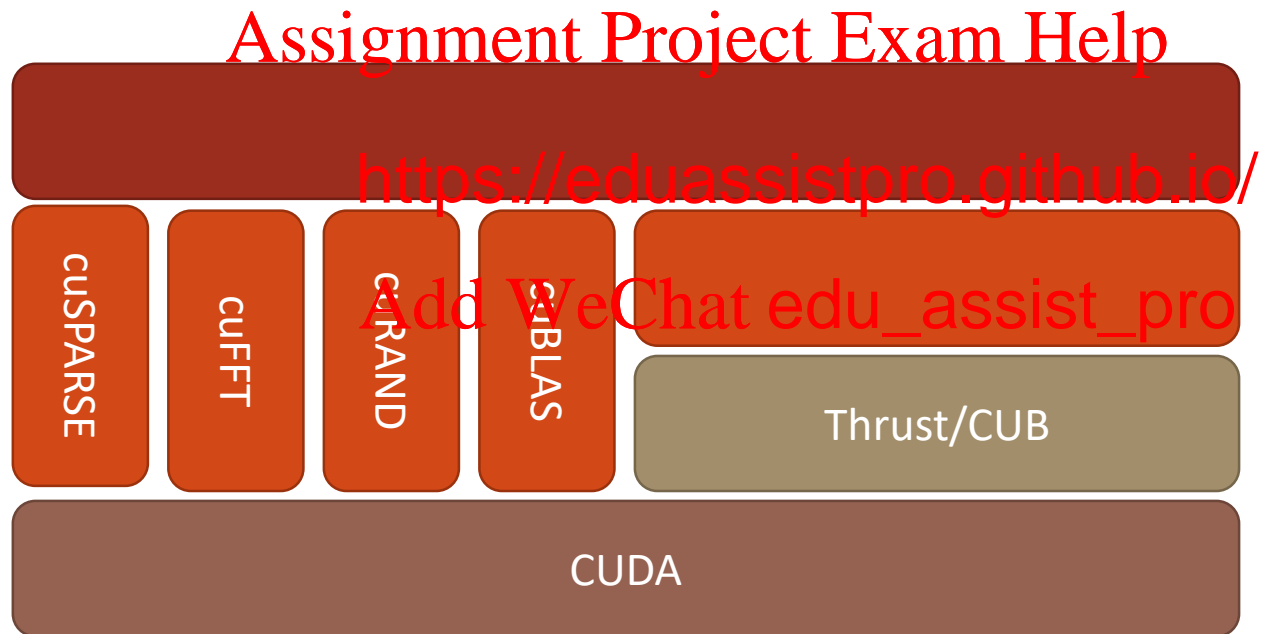
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

CUDA libraries

- ❑ Abstract CUDA model away from programmer
- ❑ Highly optimised implementations of common tools
 - ❑ Mainly focused on linear algebra



Thrust

❑ Template Library for CUDA

- ❑ Implements many parallel primitives (scan, sort, reduction etc.)
- ❑ Part of standard CUDA release
- ❑ Level of Abstraction with memcpy's



Assignment Project Exam Help

<https://eduassistpro.github.io/>

❑ Designed for C++ programmers

- ❑ Similar in design and operation as the Template Library (STL)
- ❑ Only a small amount of C++ required..

Add WeChat edu_assist_pro

Thrust containers

❑ Thrust uses only high level *vector* containers

❑ `host_vector`: on host

❑ `device_vector`: on GPU

❑ Other STL containers include

❑ `queue`

❑ `list`

❑ `tack`

❑ `queue`

❑ `priority_queue`

❑ `set`

❑ `multiset`

❑ `map`

❑ `multimap`

❑ `bitset`

❑ STL containers can be used to initialise a Thrust vector

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat `edu_assist_pro`

```
#include <thrust/host_vector.h>
#include <thrust/device_vector.h>

int main()
{
    // create a vector on the host
    thrust::host_vector<int> h_vec(10);

    // or on the device
    thrust::device_vector<int> d_vec = h_vec;

    // populated directly from host
    for (int i = 0; i < 10; i++)
        d_vec[i] = i;

    // vector memory automatically released
    return 0;
}
```

Thrust Iterators

- ❑ They point to regions of a vector
- ❑ Can be used like pointers
 - ❑ Explicit cast when dereferencing very important

Assignment Project Exam Help

```
thrust::device_vector<int> d_vec;
thrust::device_vector<int> c_vec;
printf("d_vec at begin=%d", (int)*begin);
begin++; //move on a single position
printf("d_vec at begin++=%d", (int)*begin);
*end = 88;
printf("d_vec at end=%d", (int)*end);
```

```
d_vec at begin=0
d_vec at begin++=1
d_vec at end=88
```


Thrust Iterators

❑ Can be converted to a raw pointer

```
int * d_ptr = thrust::raw_pointer_cast(begin);  
int * d_ptr = thrust::raw_pointer_cast(begin[0]);
```

Assignment Project Exam Help
kernel<BLOCKS, TPB>(d_ptr),

<https://eduassistpro.github.io/>

❑ Raw pointers can be used in Thrust [Add WeChat edu_assist_pro](#)

❑ BUT not exactly the same as a vector

```
int* d_ptr;  
cudaMalloc((void**)&d_ptr, N);  
thrust::device_ptr<int> d_vec = thrust::device_pointer_cast(d_ptr);  
//or  
thrust::device_ptr<int> d_vec = thrust::device_ptr<int>(d_ptr)  
cudaFree(d_ptr);
```

Thrust Algorithms

☐ Transformations

- ☐ Application of a function to each element within the range of a vector

☐ Reduction

- ☐ Reduction of a set of values to a single value using binary associative operator
- ☐ Can also be used to compute

☐ Prefix Sum

- ☐ Both inclusive and exclusive scans

☐ Sort

- ☐ Can sort keys or key value pairs

☐ Binary Search

- ☐ Position of a target value

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Thrust Transformations

❑ Some examples of the many transformations

```
thrust::copy(d_vec.begin(), d_vec.begin() + 10, d_vec_cpy.begin());
```

```
thrust::fill(d_vec.begin(), d_vec.begin() + 10, 0);
```

```
thrust::generate(d_vec.begin(),
```

```
//rand is a predefined Thrust ge
```

```
thrust::generate(d_vec.begin(), d_vec.begin() + 10
```

```
// fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8,
```

```
thrust::sequence(d_vec.begin(), d_vec.begin() + 10);
```

```
//all occurrences of the value 1 are replaced with the value 10
```

```
thrust::replace(d_vec.begin(), d_vec.end(), 1, 10);
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Thrust Algorithms

❑ Either in-place or to output vector

```
thrust::device_vector<int> d_vec(10);  
thrust::device_vector<int> d_vec_out(10);  
  
//fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}  
thrust::sequence(d_vec.begin(), d_vec.begin() + 10);  
  
//inclusive scan to output  
thrust::inclusive_scan(d_vec.begin(), d_vec.end(),  
    d_vec_out.begin());  
  
//inclusive scan in place  
thrust::inclusive_scan(d_vec.begin(), d_vec.end(),  
    d_vec.begin());  
  
//generate random data (actually a transformation)  
thrust::generate(d_vec.begin(), d_vec.end(), rand);  
  
//sort in place  
thrust::sort(d_vec.begin(), d_vec.end());
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Custom Transformations

```
thrust::device_vector<int> d_vec(10);
thrust::device_vector<int> d_vec_out(10);

//fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
d_vec = thrust::sequence(d_vec.begin(), d_vec.begin() + 10);
```

Assignment Project Exam Help

```
//declare a custom operator
```

```
struct add_5{
    __host__ __device__ int oper
    return a + 5;
}
};
```

```
add_5 func;
```

```
//apply custom transformation
```

```
thrust::transform(d_vec.begin(), d_vec.end(), d_vec_out.begin(), func);
```

```
//d_vec is now {5, 6, 7, 8, 9, 10, 11, 12, 13, 14}
```

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Thrust Fusion

❑ For best performance it is necessary to fuse operations

```
struct absolute{
    __host__ __device__ int operator()(int a){
        return a < 0 ? -a : a ;
    }
};
absolute func;

//custom transformation to calculate absolute value
thrust::transform(d_vec.begin(), d_vec.end(), func, result);
//apply reduction, maximum binary associate operator
int result = thrust::reduce(d_vec.begin(), d_vec.end(), 0, thrust::maximum<int>());
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
struct absolute{
    __host__ __device__ int operator()(int a){
        return a < 0 ? -a : a ;
    }
};
absolute func;

//apply transform reduction maximum binary associate operator
int result = thrust::transform_reduce(d_vec.begin(), d_vec.end(), func, 0, thrust::maximum<int>());
```

❑ Sorting Networks

❑ Merge and Bitonic sort

❑ Thrust Parallel Primitives Library

❑ Applications of sorting

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro


Sorting and parallel primitives

- ❑ Can be very useful for building data structures
 - ❑ We can use prefix sum for writing multiple values per element
- ❑ Remember Gather vs Scatter
 - ❑ What if our outputs are scattered to output
 - ❑ Very common in parti
 - ❑ Outputs might represe

Assignment Project Exam Help

<https://eduassistpro.github.io/>

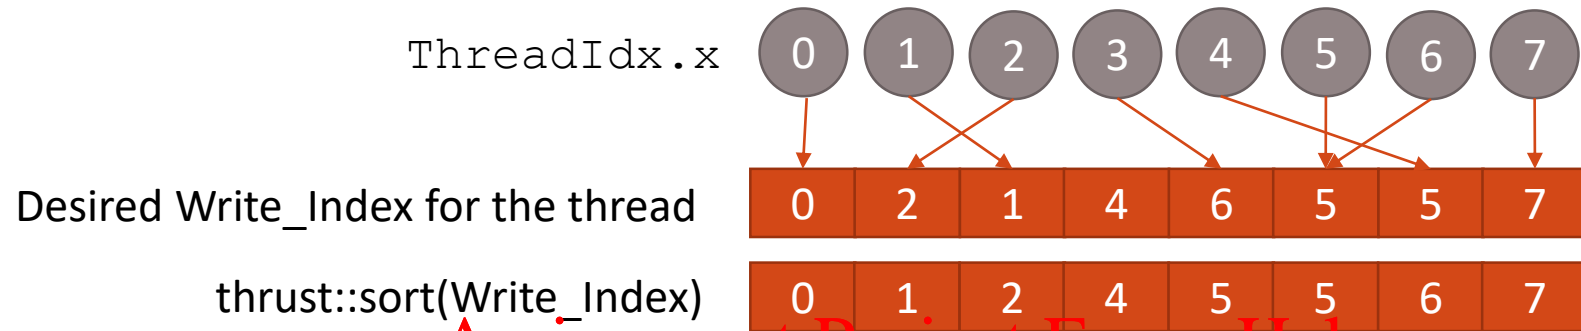
Add WeChat edu_assist_pro



Scatter operation

- ❑ Write to a number of locations
- ❑ Random access write?
- ❑ How to read multiple values afterwards?

Binning and Sorting



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Unique write indices	0	1				5	6	7
Count(Write_Index)	1	1				2	1	1
thrust::inclusive_scan(count)	0	1	2	3	3	4	6	7

← i.e. how many threads want to write to this index

❑ We can now read varying values from each bin

❑ E.g. for location 5

❑ inclusive_scan gives starting index of 4

❑ Iterate from index 4 for a count of 2 to find all values of write_index 5

Particle interaction example

- ❑ As with previous slide use sorting
 - ❑ Divide the environment according to some interaction radius
 - ❑ Output particle key value pairs (keys are location determined through some hash function)
 - ❑ Sort Keys
 - ❑ Reorder particles based on key pairs
 - ❑ Generate a partition boundary table
 - ❑ Histogram count and prefix sum
 - ❑ Each particle needs to read all particles in its own location and any neighbouring location
 - ❑ Guarantees particle interactions within the interaction radius

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

0	1	2 ¹ 2	3
3 ³ 4	4 ⁴ 5 ⁵ 6	6	7
8	9	7 ⁷ 10	11
12	13 ⁸	14	15

Partition	First agent	Last agent
0		
1		
2	1	2
3		
4	3	4
5	5	6
6		
7		
8		
9		
10	7	7
11		
12		
13	8	8
14		
15		

Summary

- ❑ Sorting networks allow data independent sort algorithms to map easily parallel architectures
- ❑ Choice of a sorting network will dictate the memory access pattern and hence the performance on a GPU
- ❑ Merge sort and Bitonic sort are good choices for GPUs
- ❑ Thrust implements many parallel primitives
- ❑ Thrust is based on the idea of containers, transformations and algorithms
- ❑ Sorting can be used to improve complex problems such as particle systems over a fixed range

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Acknowledgements and Further Reading

❑ Comparison on sorting approaches on GPU

❑ <http://arxiv.org/ftp/arxiv/papers/1511/1511.03404.pdf>

❑ <https://devblogs.nvidia.com/parallelforall/expressive-algorithmic-programming-thrust/>

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro